

# Real-Time Analytics on Data Lakes: Indexing Amazon S3 for up to 125x Faster Queries

**Dhruba Borthakur** / CTO, Rockset

**Nadine Hachouche** / Senior Developer Advocate, Rockset

Presented at Data Riders Meetup 04/2021

Slack Community: [bit.ly/rockset-community-channel](https://bit.ly/rockset-community-channel)

# Presenter



**Dhruba Borthakur**  
Co-Founder & CTO



**Nadine Hachouche**  
Senior Developer Advocate


# Agenda

1. The state of cloud data lakes
2. Rockset for real-time analytics
3. Considerations for apps on data lakes
  - a. Indexing vs scanning
  - b. High concurrency
  - c. Mutability of data and schema
4. Developer productivity
5. Workshop


# The state of cloud data lakes

# Data lake adoption continues to rise

- Store raw data as is
- Bring multiple types of data together
- Scale to massive volumes of data cost effectively
- Run many different kinds of analytics
- Democratize access to data




# Most enterprises currently use a data lake



# Cloud data lakes outpacing on-prem

Is your organization planning to augment or replace its existing BI, analytics, and data warehousing systems with any of the following systems or cloud-based services, solely or in combination? (Please select all that apply.)



Source: TDWI (2018)

Figure 4. Based on answers from 232 respondents.


# Amazon S3 is the leading cloud data lake option

- Reliable
- Practically infinite scalability
- Compatible with many other services
- Tens of thousands of data lakes on S3



# Athena commonly used for ad-hoc queries on S3

- Standard SQL
- Based on Presto
- Serverless
- Ideal for ad-hoc queries on S3 data lakes



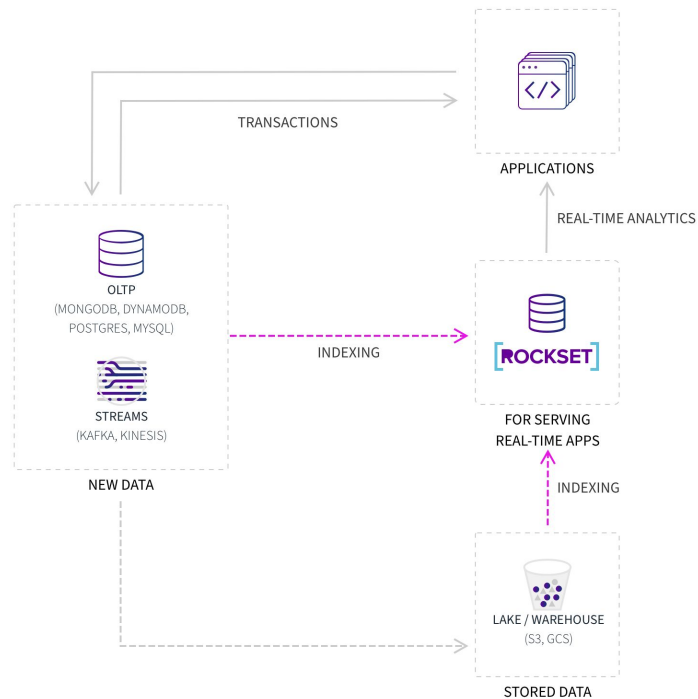
# But what about building apps on data in S3?

- Occasional queries → continuous, highly concurrent queries
- Growing need for apps that activate data in S3
  - A/B experiments on behavioral data
  - Personalization and customer 360 on marketing data
  - IoT apps on sensor data
- Need a real-time analytics solution built for low-latency, high-concurrency queries

# Rockset for real-time analytics

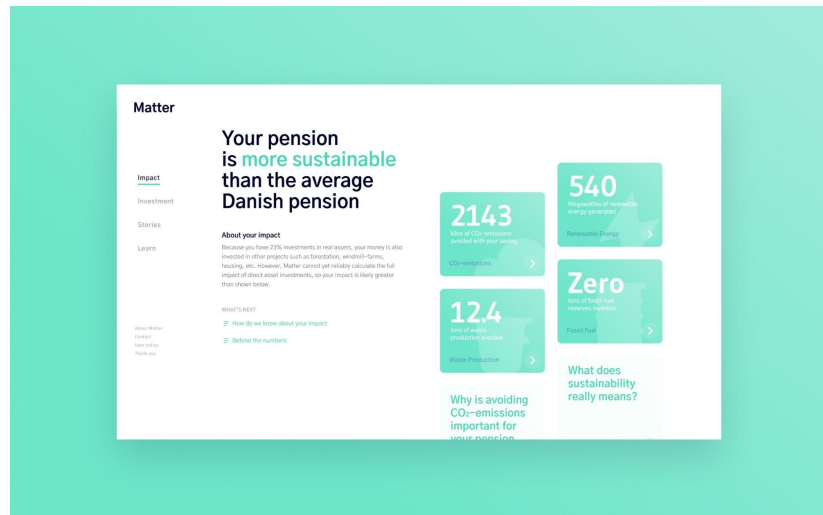
# Real-time analytics on data lakes with Rockset

- Rockset: real-time indexing database
- Built-in connector to S3
- Continuous, schemaless ingest
- Purpose built to serve apps



# Matter delivers AI-powered insights for sustainable investing

- NLP pipeline on newsfeeds stores results in S3 data lake
- High-volume queries for 1000s of asset positions in investment portfolios
- Hours → sub-second query latency moving from Athena to Rockset



# eGoGames: esports platform for mobile games

- Understand what users are doing in real time
- Analyze user acquisition and retention data in S3 with transactional data in DynamoDB
- Takes too long to centralize data in S3 to query with Athena
- Rockset allows eGoGames to query across data sources within seconds of data being produced



# Considerations for apps on data lakes


# Query Latency

Athena	Rockset
Designed for ad-hoc queries	Designed for real-time applications <ul style="list-style-type: none"><li>• high selectivity on large data sets</li><li>• continuous queries</li></ul>
Stores data in columnar format	Stores data in <ul style="list-style-type: none"><li>• rowstore</li><li>• columnstore</li><li>• inverted index</li></ul> Data sorted by <i>_event_time</i> for fast time-range queries
Every query is parallelize and scan	Queries are served from indexes <ul style="list-style-type: none"><li>• up to 125x faster than Athena</li></ul>



# Optimize Query Latency by Converged Indexing

- Traditional approach: Parallelize and scan

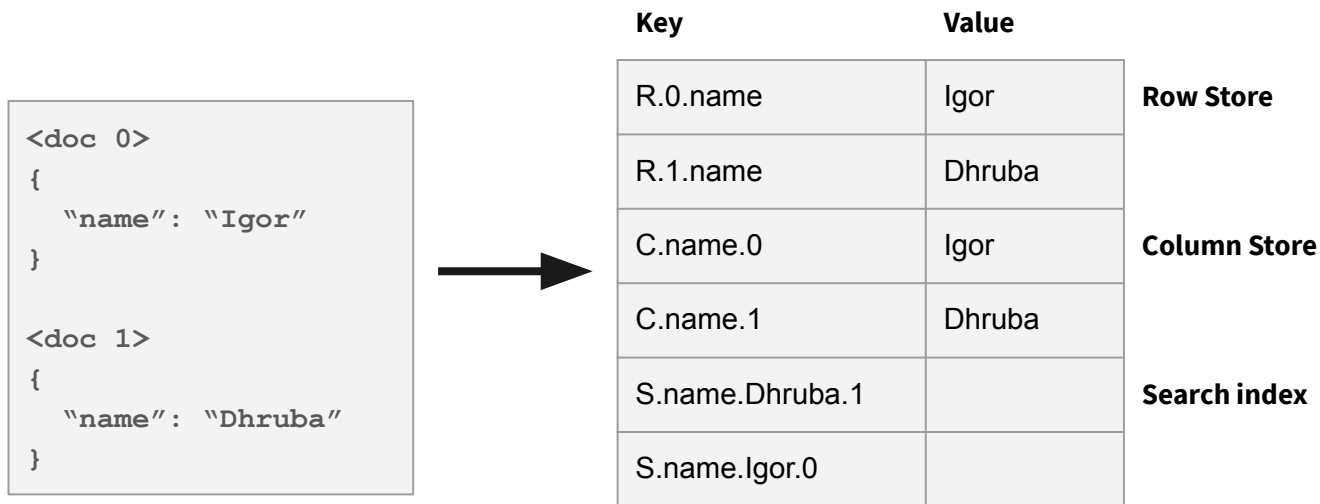


- Real-time event analytics: Parallelize and index



# Converged Indexing

- Columnar and search indexes in the same system
- Built on top of key-value store abstraction
- Each document maps to many key-value pairs



# Query Optimizer

- Low latency for both highly selective queries and large scans
- Optimizer picks between
  - inverted index (Index Filter operator)
  - columnar format (Column Scan operator)
  - inverted index (Index Scan operator)

```
SELECT *  
FROM search_logs  
WHERE keyword = 'hpts'  
AND locale = 'en'
```

**Inverted index**  
**(for highly selective queries)**

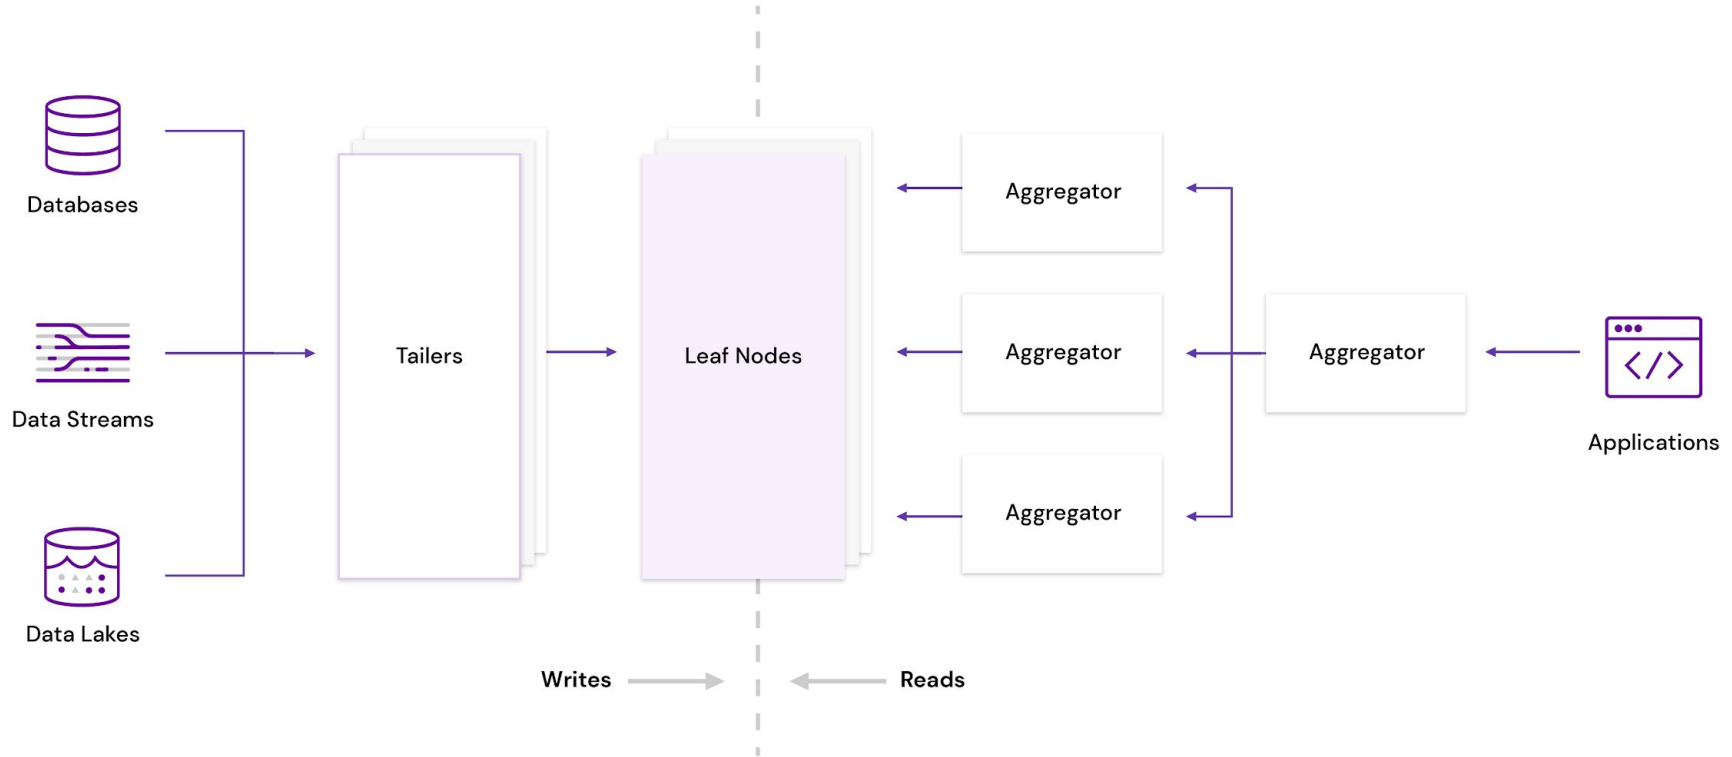
```
SELECT keyword, count(*)  
FROM search_logs  
GROUP BY keyword  
ORDER BY count(*) DESC
```

**Columnar store**  
**(for large scans)**

# Concurrency

Athena	Rockset
Used for low-concurrency use cases <ul style="list-style-type: none"><li>• ad-hoc, interactive queries</li><li>• data science</li><li>• BI</li></ul>	Used to power high-concurrency use cases <ul style="list-style-type: none"><li>• real-time analytic applications</li><li>• large numbers of users</li><li>• spiky usage</li></ul>
Executes 5 concurrent queries <ul style="list-style-type: none"><li>• queues any additional queries</li></ul>	Supports 1000s of QPS

# Rockset Uses an ALT Architecture



# Mutability of data and schema

Athena	Rockset
Requires table creation with a schema	Automatically generates schema based on exact fields and types in the data
Schema changes result in delays in querying data	Optimized for data latency <ul style="list-style-type: none"><li>raw data is immediately queryable without requiring a schema</li></ul>
Only supports inserts but not updates	All documents are mutable

# Strong Dynamic Typing

- Fields are dynamically typed

```
{"name": "Tudor", "age": 40, "zip": 94542}  
{"name": "Lisa", "age": 21, "zip": "91126"}  
{"name": "Hana"}  
{"name": "Igor", "zip": 94110.0}  
{"name": "Venkat", "age": 35, "zip": "94020"}  
{"name": "Brenda", "age": 44, "zip": "90210"}
```

# Strong Dynamic Typing

- Fields are dynamically typed
- Queries are strongly typed

```
SELECT 1 > 'a';
```



**Error [Query]**

Invalid comparison between int and string.



# Strong Dynamic Typing

- Fields are dynamically typed
- Queries are strongly typed
- Smart schemas

```
$ rock sql  
> describe tudor_example1;
```

field	occurrences	total	type
['_meta']	6	6	object
['age']	4	6	int
['name']	6	6	string
['zip']	1	6	float
['zip']	1	6	int
['zip']	3	6	string

# Set field mappings and retention as needed

▼ Field Mappings (1)

Optionally apply transformations to incoming data, such as masking sensitive fields.

**Default Document Settings**

☒ Keep all fields, not explicitly dropped.

☐ Drop all fields, not explicitly created from Field Mappings to whitelist specific fields.

Field Mapping Type

Drop Field

Choose a Field Mapping Type.

Input Field Name in Original Source

field

Required: Find the field name in your source data.

Drop Field

True

⊖ Remove

+ Add Field Mapping

▼ Retention

If specified, Rockset will drop documents after a given duration using the `_event_time` field to determine document age. By default, Rockset will use insertion time into the collection as `_event_time`. You can also map a field in your data to `_event_time` using a Field Mappings.

☒ Keep all documents


☐ Drop documents after 

30

days


# Data Latency

- Fast ingestion
  - New data is visible in query results within a minute or less
  - Continuous Live sync of new data from S3 to Rockset



# 1-sec data latency when writing 1B documents/day

4XLarge



# Developer productivity

# Query Lambdas

- Named, parameterized SQL queries stored in Rockset
- Executed from dedicated REST endpoint
- Organize by versions and tags
- Create data APIs used by multiple application developers
- Avoid having SQL in application code

:// Curl

> CLI

JS NodeJS

Python

```
curl --request POST \  
--url https://api.rs2.usw2.rockset.com/v1/orgs/self/ws/commons/lambdas/rankStocks/versions/815ef07a1fe75692  
-H 'Authorization: ApiKey [REDACTED]' \  
-H 'Content-Type: application/json' \  
| python -m json.tool
```

# Workshop

# Using Rockset to Build Real-time Analytics



## STEP 1

**Create an account and login to Rockset**

Rockset is a fully managed cloud service



## STEP 2

**Connect to your data source**

Rockset builds Converged Indexes™ for you



## STEP 3

**Save your SQL statement as a Query Lambda**

You get a REST endpoint for your data API



## STEP 4

**Hit the REST endpoint from your application code**

Get results in milliseconds



# Indexing S3 Using Rockset

- Low-latency, high-concurrency queries for serving applications
- Increase developer velocity and reduce time to market
- Complete solution for real-time analytics on data lakes

“While you can build or buy solutions individually that might provide more ingest options, better absolute performance, lower marginal costs or higher scalability potential, I have yet to find anything that comes remotely close to Rockset on all of these areas at the same time, in a setting where time to market is a highly valuable metric.”

- Matter CTO Alex Harrington



Get started with **\$300** in free credits

<https://console.rockset.com/create>

# Thank you

Dhruba Borthakur / [\*\*dhruba@rockset.com\*\*](mailto:dhruba@rockset.com)

Join the community channel after the talk at  
[bit.ly/rockset-community-channel](https://bit.ly/rockset-community-channel)